

Problem 2: Logistic Regression

Data Manipulation

I've experimented with three approaches:

1. Removing **all** missing information and highly correlated attributes.
2. Removing **all** missing information.
3. Using linear regression to fill in **some** missing information.

The following was implemented in all three approaches:

The dataset contains both *NULL* values and *0*'s where information is missing. The *0*'s have been converted to *NULL* for missing-data consistency ⁽⁵⁾.

Most of the missing information is in the *SkinThickness* and *Insulin* columns ^(6 - missingno).

There are 44 rows that contain *NULL* values **outside** of those two columns ⁽⁷⁾. These rows compose about 5.7% of the dataset and will be dropped ⁽¹¹⁾.

The *Insulin* column shows high variability (large number of outliers ^(9 - boxplot)) and a high correlation (0.58) with the *Glucose* column ^(10 - corrMatrix). Because of this, it is dropped ⁽¹²⁾.



Approaches #1 & #2:

The remaining missing data is limited to the *SkinThickness* column, which has a high correlation (0.65) with the *BMI* column ^(10 - corrMatrix). Because of this, it is also dropped ⁽¹²⁾.

Approach #1 drops the *Pregnancies* column due to its high correlation (0.54) with the *Age* column ^(10 - corrMatrix). I chose to drop *Pregnancies* instead of *Age* because *Age* has a higher correlation with *Outcome*.

Approach #3:

This approach uses linear regression to fill in the missing data in the *SkinThickness* column using the data in the *BMI* column, with which it shares a high correlation. This had the expected effect of increasing the correlation even more, which will probably have a negative impact on the final result.



Testing ⁽¹⁵⁾:

All three approaches were run through a logistic regression model ⁽¹⁵⁾. First, they were fitted and tested against their entire respective datasets (no train/test split) to see the overall accuracy of the logistic regression model. They were then tested using a range of train/test split ratios (50 iterations per split averaged together) to see how each affected the perceived accuracy of each model.

Results ⁽¹⁵⁾:

Having fewer highly correlated attributes in the final model improves the model's overall accuracy, even if only slightly as my models show ($0.779 > 0.776 > 0.769$). Additionally, the more valid data that's used to train the logistic regression model, the more accurate the model becomes – as seen by the change in accuracy between different train/test ratios. However, using a high train/test ratio may inflate the model's score due to the relatively small test pool. A formal proof would be needed, but I believe training a regression model using all rows of a dataset and then testing the model using the same data (as I did before doing the splits) is a better approach to gauging the overall accuracy of that model.

As a sidenote/afterthought, the linear regression models of this dataset were much better at classifying non-diabetic entries than they were at classifying diabetic entries, with a false positive rate of about 12% and a false negative rate of about 42%

(15 – confMatrix)

		Predicted Class	
		Non-Diabetic	Diabetic
Actual Class	Non-Diabetic	420	55
	Diabetic	105	144

Confusion Matrix for an instance of Model #1

Score Results

#1

0.7790

Train/Test	Score
90/10	0.7940
80/20	0.7731
70/30	0.7748
60/40	0.7677
50/50	0.7752
40/60	0.7677
30/70	0.7639
20/80	0.7627
10/90	0.7554

#2

0.7762

Train/Test	Score
90/10	0.7825
80/20	0.7756
70/30	0.7659
60/40	0.7679
50/50	0.7674
40/60	0.7695
30/70	0.7668
20/80	0.7630
10/90	0.7540

#3

0.7693

Train/Test	Score
90/10	0.7756
80/20	0.7668
70/30	0.7662
60/40	0.7669
50/50	0.7652
40/60	0.7628
30/70	0.7623
20/80	0.7619
10/90	0.7474